

IESTI01 - Lab 3

Project: KWS / Sound Classification

Introduction

In the class, we went through the steps of creating a speech recognition (or keyword spotting) system (in this case, using “UNIFEI and “IESTI”). The same basic steps can be used to create a model that recognizes and classifies other sounds, such as broken glass, toss, etc.

In this lab, you should **at least** create a KWS project using **three classes**:

- First word of your choice
- The second word of your choice.
- Silence (Try with some background noise)

A Fourth class can be added for **Unknown** words. You can use the KWS dataset “Speech Commands” (but remember that this dataset is in English) or Common Voice (Portuguese dataset).

Optionally you can try **a second and separated project**, starting with a simple, pre-made dataset, adding your target sound, and building a classifier for those sounds. Try to find sounds made with something other than your voice!

Required Hardware

For collecting sound data, you should have access to a recording device. This can be an Arduino 33 BLE Sense (TinyML Kit), smartphone, webcam, laptop, etc.

For deploying, you must use the [Arduino Nano 33 BLE Sense](#).

Collect Data for optional Sound Project (not voice)

You can test a pre-made dataset from Edge Impulse that includes a generic noise category and the sound of a faucet running. You can use that data, augment it with your sounds, or collect your own entirely new dataset.

To start, download the faucet dataset from [this link](#). Unzip it somewhere on your computer.

If you use your sounds, collect at least 50 1-second audio samples of each class (or at least one 50-second recording). Ensure the sound source is in different environments and distances away from the recording device. This will help create a more robust model to differentiate that sound from other noises.

For example, recorded a fan, a blender, or other machines running at different speeds. Move the microphone to several angles to collect the sound in front of, behind, and above the device as the sound changes as the angles and distance vary.

The idea is to have the embedded system identify at least three different types of sounds: background noise, faucet running, or machine working (select only one: fan, blender, car engine, etc.).

Curate Data

When you're done, transfer the sounds over to your computer. Remember that if you collect directly with your phone, you should have an app that samples sound using 16KHz/16bits. If not, you can use [Audacity](#) (or Librosa Library in Python to convert the sound).

Using Audacity App:

Open your audio file in Audacity. Select the drop-down menu, *Project Rate (Hz)*, at the bottom of the screen. Change the sampling rate to **16000 Hz**. And click **File > Export > Export Selected Audio**. Set the bit depth to **Signed 16-bit PCM**. Save the audio selection as a WAV file. You can use Audacity to split the sound in a 1-second clip, but using the "Split Feature" available at Edge Impulse Studio is more accessible (in my opinion).

Use the naming convention of *<label>.xxxx.wav*. The *<label>* before the first period (.) can be used by Edge Impulse to determine the class automatically. A number (XXXX) is used to identify the files uniquely. The ordering and the exact number are unnecessary (i.e., you could also use a hash).

NOTE: You can also collect audio data straight from your Edge Impulse project as we did in class! Go to *Data acquisition* in a new project and connect your smartphone or Arduino board.

Upload Data (when you have .wav data collected):

- Create a new project in Edge Impulse.
- Head to the **Data Acquisition** page.
- Click **Let's collect some data**.

- Select the **Go to the uploader** option.

On this new page, click **Choose files**. Select all of the files from your curated sample set.

- Click **Open**.

Leave *Automatically split between training and testing* selected. If you used the file naming scheme I outlined above, leave *Infer from the filename chosen*. If not, select *Enter label* and give your samples a label (e.g., “fan”).

- Click **Begin upload**.

Repeat this process for the *faucet* and *noise* sets you downloaded in the *faucet_dataset* ZIP file you downloaded at the beginning. The WAV files in that set should follow the naming scheme outlined, letting you leave *Infer from the filename* selected to add labels to the samples.

- Click on the **Data acquisition** link to go back to the Data Acquisition page. Here, make sure that all of your samples are present and that they are divided between the training and test sets (there should be about 20% of the pieces in the test set).

Feature Extraction

Navigate to the **Impulse design** page of your project.

- Add an **Audio (MFCC)** processing block for KWS (human words) or **Audio (MFE)** (for non-voice audio)
- Add a **Neural Network (Keras)** learning block.
- Click **Save Impulse**.

Go to the **Spectrogram** page.

- Click on the **Generate Features** tab.
- Click the **Generate Features** button, and wait a moment while your audio samples are converted into spectrograms. When it's done, look at the *Feature explorer* to see if you can identify separation among your classes.

Model Training

Navigate to the **NN Classifier** page. Please leave all of the hyperparameters at their defaults and click **Start training**. When it's done, scroll down to view the *Confusion matrix* of the validation data.

- Download the Jupyter Notebook (Expert Mode) and run it at CoLab.
- Create the history graphs, try changing some of the hyperparameters, and re-training your model to see if it improves the per-class accuracy.
- Return to Edge Impulse Studio and re-train your model with eventual different hyperparameters.

Testing

Head to the **Model testing** page and classify all of the test data. Eventually, you can try Live Test using the Kit. If you're happy with the test results, continue to the deployment step (I hope to see an accuracy better than 75%, if not, go back to collect more data and adjust hyperparameters).

Deployment

Head to the **Deployment** page in your Edge Impulse project.

- Select the **Arduino library** and click **Build**.
- When it's done downloading, open your Arduino IDE.
- Select **Sketch > Include Library > Add .ZIP Library...** Select the library file you just downloaded.
- Go to **File > Examples > <your-project-name> > nano_ble33_sense_microphone**
- Click **Upload** to compile and send the program to your Arduino board.

When it's done, open a Serial Monitor. Try holding the Kit (Arduino board) up to different noise sources.

Arduino LEDs

Modify the Arduino sketch **nano_ble33_sense_microphone** to include the RGB LEDs (do it in a way that you can use the "stand-alone" w/o looking at IDE Serial Monitor to see what class your project is selecting).

Report

For grade, delivery a report with relevant photos/print screens and the CoLab with your project. Optionally you can provide videos.